

# HW1 - High Performance Computing

Andrew Szymczak

February, 2015

## GitHub Repository



**MPI ring communication.** Write a distributed memory program that sends an integer in a ring starting from process 0 to 1 to 2 (and so on). The last process sends the message back to process 0

- I allow two command line parameters  $N$  and  $M$ .  $N$  specifies the number of loops and  $M$  specifies the number of integers sent in each message. Each message is an array of  $M$  integers.

```
msg = (int *) malloc(sizeof(int) * M);
```

The first element is the only one altered as `msg` is passed around; `msg[0] += rank`. If you don't input an argument or you input an illegal value,  $N, M$  default to 1. On my machine `ints` are 4 bytes, so to send 2MB messages I input  $M = 524288$ . If the number of processors is less than 1, the program aborts.

- After the last message is received at processor zero, I verify that `msg[0] = NP(P-1)/2` where  $P$  is the number of processors.
- I only tested my program locally. The latency for a message seems to be about .026 ms. 20 processors were used to make the following plots.

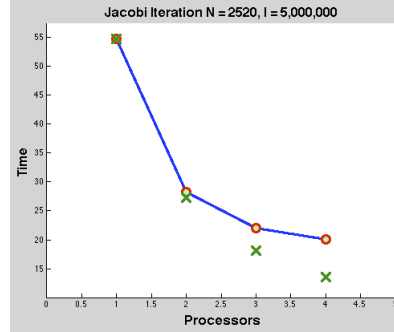


- With  $P = 20$ ,  $N = 10000$  and  $M = 524288$ , a total of 200000 messages were sent in 124.71 seconds. Since the limiting factor is bandwidth rather than latency, I can estimate a bandwidth of about 3.13 GB/s.

**Distributed memory parallel Jacobi smoother. Use MPI to write a parallel version of the Jacobi smoother from Homework 0**

- The first and last element of  $\mathbf{u}^{(i)}$  are the shared indices. First we pass the previous values  $u_{-2}^{(i)} \rightarrow u_0^{(i+1)}$  for all  $i < p-1$ . Then we perform the Jacobi iteration. Then we pass the new values  $u_1^{(i)} \rightarrow u_{-1}^{(i+1)}$  for all  $i > 0$ .
- I use an `MPI_Allreduce` to calculate the residual on every processor. Since this is an expensive process, I only do it once every 100 iterations.
- Since my machine only utilizes four processors, my first scaling test was for  $p = 1, 2, 3, 4$ . Let  $I$  denote the number of iterations, and  $t$  the total time in seconds.

$N$	$I$	$p$	$ Au - f $	$t$
2520	5000000	1	0.932	54.62
2520	5000000	2	0.932	28.26
2520	5000000	3	0.932	22.08
2520	5000000	4	0.932	20.12



Note how the residual is independent of  $p$ , as expected. Each processor handles 2 messages and  $N/p$  of the Jacobi work in each iteration. Assuming full connectivity (so that all messages are done in parallel), we can expect  $\mathcal{O}(N/p)$  time (as marked by the X's). The slowdown is due to the suboptimal parallelisation and overhead of message passing.

- My second scaling test only establishes the high overhead of message passing. Increasing  $p > 4$  on my local machine effectively sequentializes crosstalk. Since there are  $2p-2$  messages passed in every iteration, we can expect the computation time to grow  $\sim p$ .

$N$	$I$	$p$	$ Au - f $	$t$
1200	100000	20	22.24	10.40
1200	100000	40	22.24	21.83
1200	100000	60	22.24	33.17
1200	100000	80	22.24	48.06
1200	100000	100	22.24	63.77

- Gauss-Seidel would be more difficult to program because of the implicit dependence in each iteration. We could try to do something analagous to the first bullet point: First we pass the old values  $u_1^{(i)} \rightarrow u_{-1}^{(i+1)}$  for all  $i < p-1$ . Then we perform the Jacobi iteration. Then we pass the new values  $u_{-2}^{(i)} \rightarrow u_0^{(i+1)}$  for all  $i < p-1$ . The problem is that you can't do this in parallel. Before  $u^{(i)}$  can preform Jacobi, it must wait to recieve the new value  $u_{-2}^{(i-1)}$ , which isn't sent until all lower rank processes are complete. The process is essentially sequential, and even worse off due to message passing.